



# The Prompt Engineer's Compendium

Complete Guide to Expert Prompting Across All Major AI Platforms

## Section I — Foundations of Expert Prompting

### Unifying Principles

#### 1. Clarity, Specificity, Context

**Core Truth:** Vague prompts  $\Rightarrow$  generic outputs

Always add: audience, length, format, tone, and constraints to your prompts.

**✗ Bad:** "Find a restaurant"

**✓ Good:** "Best Indian restaurant in Cambridge, Massachusetts, within walking distance of Harvard Yard."

#### 2. Affirmative Directives (Tell the model what TO do)

Prefer positive instructions over negatives for better comprehension.

**✓ Better:** "Diagnose the issue while refraining from any PII-related questions."

**✗ Avoid:** "Do not ask for username or password."

### 3. Deliberate Structure

Use clear delimiters and sections so the model can parse reliably.

**Common patterns:** ### Headings, fenced blocks ("""), bullet lists, and tagged segments.

#### Platform Nuances:

- **Claude:** Favors XML-like tags for high-fidelity parsing (e.g., <instructions>...</instructions>, <context>...</context>, <examples>...</examples>)
- **Grok:** Responds well to Markdown structure and concise sections
- **Tip:** For complex tasks, separate rules, context, examples, and output format with clear labels

## Section II – Prompting Paradigms

### 1 Chain-of-Thought (CoT)

Guides the model to show intermediate reasoning before the final answer.

**Trigger phrases:** "Let's think step by step.", "First, let's reason this out.", "Work it out step by step to ensure the right answer."

**Classic demonstration (the "Apple Problem"):** A direct ask often fails; appending CoT prompts the model to enumerate steps (e.g.,  $10 - 2 - 2 + 5 - 1$ ) and reach the correct result.

**Use for:** Math, logic, planning, and multi-constraint synthesis.

```
### CoT Template
### Task
[SPECIFIC QUESTION]
```

```
### Method
Let's think step by step. Show intermediate
calculations/reasons before the final answer.

### Answer
```

## 2 Shot-Based Instruction (In-Context Learning)

**Zero-Shot:** Just the instruction. Good for simple, common tasks.

**One-Shot:** Add one example to set tone/format for slightly nuanced tasks.

**Few-Shot (2+ examples):** Best for accuracy and consistency; teaches pattern/format.

```
### Few-Shot Skeleton
### Instructions
[WHAT TO DO]

### Examples
Input: ...
Output: ...

Input: ...
Output: ...

### Now Do This
Input: [YOUR INPUT]
Output:
```

## 3 Persona & Role-Playing (Use with care)

Assign a role to influence tone/focus (e.g., "You are an expert sentiment classifier.").

**Caveat:** Misaligned personas can degrade factual or reasoning performance (e.g., irrelevant role primes the wrong context).

**Mitigation — "Jekyll & Hyde" ensemble:** Generate two answers (persona vs neutral) and select the better one via an evaluator prompt.

```
### Persona + Neutral Ensemble Pattern
### Version A (Persona)
<role>Senior [DOMAIN] Expert</role>
[instructions]

### Version B (Neutral)
[instructions, no role]

### Judge
You are the evaluator. Compare Answer A (persona) vs Answer B
(neutral) for correctness, completeness, and adherence to
constraints.
Return ONLY:
{
  "winner": "A|B",
  "reason": "<1-2 sentences>"
}
```

## 4

## Advanced Reasoning Patterns

Use these when you need super-reliable reasoning, planning, or tool use.

**Platforms:** ChatGPT, Claude, Grok, Perplexity (research), Gemini (planning), NIM (as endpoint in ReAct).

## Advanced Reasoning Pattern Templates

## A) ReAct (Reason + Act)

```
You are a planner-agent that alternates THOUGHT and ACTION to solve the
task. Tools available: [search()], [calculator()], [code_runner()],
[db.query()]. Task: [clear goal] Rules: - Always write a THOUGHT before
any ACTION. - After each ACTION, record OBSERVATION. - Stop when you
have enough evidence; return FINAL ANSWER with citations/derivations.
THOUGHT: ... ACTION: search("...") OBSERVATION: ... THOUGHT: ... ACTION:
calculator("...") OBSERVATION: ... FINAL ANSWER: ...
```

## B) Tree-of-Thought (ToT)

```
Goal: [problem]. Generate 3 distinct solution paths (Thought A/B/C). For
each path: list steps, assumptions, risks. Score each on (correctness,
cost, coverage) 1-5. Pick the highest-scoring path and execute it step-
by-step before concluding. Return: (1) short rationale of why the chosen
path wins, (2) final answer.
```

## C) Self-Ask (Decompose → Solve → Compose)

```
Decompose the query into sub-questions (SQ1..SQn). Answer each
succinctly. Then synthesize a final answer that cites which sub-answers
support each claim. Query: [your question] Output sections: Sub-
Questions • Answers • Synthesis.
```

## D) Graph-of-Thoughts (GoT)

Transform the task into a graph of micro-tasks. Nodes: [N1..Nn] with purpose, inputs, outputs. Edges: dependencies. Execute nodes in topological order; allow back-edges for refinement. Return: the graph (adjacency list) + final result.

## E) Self-Consistency Voting

Produce 5 independent reasoning paths (no reuse). Return just the final answers for each, then select the winner via a brief majority vote + 1-2 line justification.

## F) Chain-of-Verification (CoVe)

Task: [claim-heavy task]. 1) Draft: Produce a concise initial answer. 2) Verify-Plan: List 5-8 targeted verification questions. 3) Verify-Answer: Answer each verification question independently. 4) Revise: Update the draft, explicitly noting changes. 5) Confidence: 0-1 with a one-line caveat. Output sections: Draft • Verification Qs • Answers • Revised Answer • Confidence.

## G) Reflexion (Critique → Improve Loop)

Round 1 – Answerer: Produce your best answer to [task]. Round 2 – Critic: List concrete flaws, missing cases, weak assumptions, and style issues. Round 3 – Answerer: Revise to address every critique point. Mark

changes with ►. Stop after one loop unless quality < A-; if so, run one more loop. Return: Final answer + bullet list of fixes applied.

### General Task Card Template

```
### Role <role>[Optional persona aligned to task]</role> ### Goal [One sentence objective] ### Context [Audience, tone, length, format, constraints] ### Examples (Optional Few-Shot) Input: ... Output: ... ### Output Format - Return as: [specify format]
```

## Section III – Platform Strategy (Strengths & Styles)

Platform	Key Strengths	Unique Use Cases	Prompting Style Highlights
ChatGPT	Versatile, general-purpose; strong for text/code and multi-task workflows	Creative writing, prototyping, structured extraction	Clear sections, explicit formats, examples/delimiters
Gemini	Multimodal + real-time; deep Google ecosystem	Photo-to-figurine (Nano Banana), image/video tasks	Structured, multi-component prompts combining text + images
Claude	Strong reasoning; massive context; safety	Long-form analysis, legal/academic, structured data	XML-tagged segments for instructions/context/examples
Grok	Fast, conversational; access to X (formerly Twitter) data	Real-time threads/posts;	Concise Markdown, iterative refinement, natural language

Platform	Key Strengths	Unique Use Cases	Prompting Style Highlights
		motion-oriented content	
Perplexity	Research-centric with citations	Up-to-date synthesis, comparative reviews	Pose as research queries with tight scope/keywords
NVIDIA NIM	Optimized inference microservices (API-first)	3D asset generation pipelines; enterprise apps	API-driven prompts inside structured requests

### Match Tool to Task

- General creativity: ChatGPT or Gemini
- Long-form analysis: Claude
- Live research: Perplexity
- Fast conversational: Grok
- 3D assets & enterprise: NVIDIA NIM

### Design Multi-Stage Systems

Consider persona+neutral ensembles ("Jekyll & Hyde") with an evaluator step for balanced outputs

### Structure + Examples

Delimiters and few-shot patterns reduce ambiguity and improve fidelity across all platforms

# Section V – Copy-Paste Prompt Library (By Task)

## A) Text-to-Website & Code Generation

End-to-end scaffold + code



*Purpose: Complete website architecture and implementation*

Develop an architecture and code for a [e-commerce] website with JavaScript. Specify pages, components, state, and data flow; generate modular code.

ChatGPT

Gemini

Claude

Grok

## Implement specific web feature

*Purpose: Create sticky header with best practices*

Act like a programming languages expert. I need a sticky header. Provide a CSS + JS example and explain pitfalls (z-index, scroll throttling).

ChatGPT

Gemini

Grok

Claude

## Debug and harden code

*Purpose: Find bugs and provide fixes*

Find the bug in this code: [paste code]  
Explain cause and provide a minimal fix + test.

ChatGPT

Claude

Grok

Gemini

## Output Format Hint

```
### Constraints
- Use semantic HTML, ARIA labels, and responsive design.
- Return code blocks per file: /index.html, /styles.css, /app.js.
- Include a 60s Lighthouse checklist.
```

## B) Social Media Content & Marketing

### High-engagement X thread

*Purpose: Create viral Twitter/X thread*

Write a 7-tweet thread on [topic] that builds curiosity, delivers real value, and ends with a CTA to follow. Include hooks, numbered tips, and a summarizing last tweet.

Grok

ChatGPT

Gemini

## LinkedIn post from voice notes

*Purpose: Transform raw notes into professional post*

Turn this voice-note transcript into a high-authority LinkedIn post—punchy, no fluff, leave space for engagement. [paste text]

ChatGPT

Gemini

Claude

## Multi-channel content plan

*Purpose: 3-month marketing strategy*

I'm launching a new [product]. Create a 3-month content marketing plan across blog, social, and email. Include cadence, topics, and CTAs.

ChatGPT

Gemini

## SEM planning

*Purpose: Keyword research and ad strategy*

Act as a digital marketing manager for [company]. Seed keywords: [list]. Expand to keyword clusters (intent, volume guess), ad copy angles, and negative keywords.

ChatGPT

Gemini

Perplexity

## C) Presentations & Pitch Decks

## Structured pitch deck plan

*Purpose: Complete investor presentation outline*

Create a presentation outline and draft content for a [10-minute pitch deck] about [product] for an [investor] audience. Main message: XX. Key claims: 1-3. Include slide titles, bullets, and suggested visuals.

ChatGPT

Claude

## Slide-level refinement

*Purpose: Polish individual slides*

From the outline above, write 2-5 sentences that support this slide's main message. Suggest visuals and a transition to the next slide.

ChatGPT

Claude

Gemini

## Slide Output Format

Return as:  
Slide #: Title  
• Key point 1  
• Key point 2  
Visual suggestion: ...  
Transition: ...

## D) Visual & Multimedia Creation

### D1) Photo → Figurine (Nano Banana)

*Purpose: Transform photo into commercial figurine render*

Create a 1/7 scale commercialized figurine of the characters in the picture, in a realistic style, in a real environment.  
The figurine is placed on a computer desk.

The figurine has a round transparent acrylic base, with no text on the base.

The content on the computer screen is a 3D modeling process of this figurine.

Include a toy packaging box in the background that matches the figurine.

Keep lighting natural and reflections realistic.

Gemini

## D2) Short 3D Animated Scene

*Purpose: Create animated video clip*

[Style]: a short 3D animated scene in a joyful cartoon style.

[Subject]: A cute creature with snow-leopard-like fur.

[Action]: happily prances through a whimsical winter forest.

[Context]: rounded, snow-covered trees; gentle falling snowflakes; soft ambient light.

[Extras]: simple SFX for footsteps; end on a playful leap freeze-frame.

Duration: ~8-12 seconds.

Gemini (Veo)

Grok

## E) Product Deep Research

### Comparative Research

*Purpose: In-depth product comparison*

Pro Search: Compare [Product A] vs [Product B] on features, pros/cons, price.

Scope: focus on [industry/region]. Return a table and 5+ cited sources. Include a 2-sentence bottom line.

Follow-ups:

- Narrow scope to sources after: 2024-01-01
- Include only primary sources and whitepapers; exclude affiliate blogs
- What are the strongest counterarguments? Cite at least 3
- Summarize consensus vs open disagreements in 5 bullets

## Research Brief

*Purpose: Structured research summary*

Labs: Build a research brief on [topic]. Sections: Background, Key Findings, Risks/Open Questions, Further Reading (linked). Limit to 400 words; cite sources inline.

## Long-Context Summarization

*Purpose: Hierarchical summary with citations*

```
<instructions>
  Create a hierarchical summary with section headings and bullet
  points.
  Append an APA-style mini-bibliography using the provided refs.
</instructions>
<sources>
  [DOC 1]
  [DOC 2]
</sources>
<constraints>
  400-600 words; keep quotations under 20 words each.
</constraints>
```

## F-J) Creative & Advanced Prompts

### F) Animated Video + SFX

*Purpose: Motion graphics with sound*

Create a 10-second animated clip: a paper airplane weaving through neon city canyons at night.  
Add gentle whoosh SFX synced to turns; end with a soft chime.  
Style: playful, minimal; include a witty caption and 3 on-trend

hashtags.  
Deliverables: MP4 + caption text.

Grok

## G) Worldbuilding Bible

*Purpose: Complete fictional world design*

Create a worldbuilding bible for a [genre] mini-series.  
Sections: Premise (2 lines), World Rules (7 bullets), Factions (3), Key Locales (5), Core Cast (6 with desires & secrets), Season Arc (8 beats), Visual Lookbook (10 image prompts), Soundtrack (moods + instruments), Iconic Props (7).  
Style: concise, high-signal, production-ready.

ChatGPT

Claude

## H) Brand Voice Cloner

*Purpose: Extract and replicate writing style*

Given 3 writing samples (below), infer a Style Guide with: Tone sliders, Diction, Syntax patterns, Tropes to use/avoid, CTA patterns, and Do/Don't examples.  
Then rewrite this draft in the inferred voice with a 1-paragraph rationale.  
[SAMPLES]  
[DRAFT]

ChatGPT

Claude

## I) Idea Oracle

*Purpose: Generate ideas through multiple lenses*

Topic: [space]. Generate 12 ideas via lenses: (1) First principles, (2) Inversion, (3) Constraints-as-features, (4) Analogy, (5) Edge cases, (6) Opposite day.  
Return a 2-column table: Idea • Why it could work.

ChatGPT

Claude

Gemini

## J) Cinematic Shot-List Factory

*Purpose: Film production planning*

Turn this scene idea into a 30-shot list with shot type, lens, movement, and motivation. Tag B-roll and cutaways. Add 3 alt shots for safety.

Scene: [description]

ChatGPT

Gemini

## Additional Creative Prompts

### Fusion Style Writer

*Purpose: Blend writing styles*

Write [format] that fuses [Author/Style A] + [Author/Style B] while avoiding direct mimicry; describe the blend in 2 lines first.

Topic: [topic]. Length: [x].

ChatGPT

Claude

### "Director's Prompt" Formula (Image)

*Purpose: Detailed image generation*

[Subject] • [Action] • [Environment] • [Composition] • [Lighting] • [Camera/lens] • [Color/mood] • [Materials/textures] • [Post-processing].

Example: A tiny clockwork dragon • landing on a tea cup • on a sunlit windowsill • rule-of-thirds close-up • rim-lit with soft bloom • 85mm shallow DoF • warm amber • brass + porcelain • subtle film grain.

Gemini

ChatGPT

### Nano-Banana Variants

*Purpose: Different figurine styles*

Make a 1/7-scale commercial figurine of the subject from the photo in [style]: (diorama / papercraft / LEGO brick-built / claymation). On a tidy desk; clear acrylic round base; no text; natural reflections.

Gemini

## **Veo Short – Action + SFX**

*Purpose: Short video with sound effects*

Create an 8-12s clip: [subject] performs [action] in [setting]. Include [SFX beats] synced to motion; end on a logo sting. Output MP4 + caption.

Gemini (Veo)

## **Hook Factory (20 Frameworks)**

*Purpose: Generate multiple hook variations*

Generate 20 hooks for [topic] using frameworks: AIDA, PAS, ABT, 4U, Before-After-Bridge, Problem-Myth-Truth, Data-Gap-Payoff, Numbered-List, Hot-Take, Contrarian, Mini-Story, Question, Challenge, Stat-Shock, Template-Madlib, Insider-Tip, Mistake-to-Avoid, Debunk, Case-Study, Objection-Smash. Return a table: Hook • Framework • Angle.

ChatGPT

Claude

## **Research Thread Distiller**

*Purpose: Convert research into social media thread*

Turn these sources into a 7-post X thread with a bold claim, 5 supported points (1 link each), and a takeaway + CTA. Include an alt-text line per image. [links]

Grok

ChatGPT

## **Meme Engine**



*Purpose: Generate meme concepts*

Propose 5 meme concepts for [topic]. For each: template name, caption text, alt text for accessibility, and why it will resonate with [audience].

ChatGPT

Grok

## Technical & Code Analysis Prompts

## Property-Based Test Generator

*Purpose: Generate comprehensive test cases*

Given code [paste], infer properties/invariants and generate property-based tests (e.g., Hypothesis/QuickCheck). Include edge cases and shrinking hints.

ChatGPT

Claude

## Targeted Verification Qs for Code

*Purpose: Pre-merge code review*

For this diff, generate verification questions that would uncover hidden bugs (nullability, bounds, concurrency, perf). Map each Q to the exact line(s) touched. [paste diff]

Claude

ChatGPT

## Repo Mental Model

*Purpose: Understand codebase architecture*

Create a bird's-eye map of this repo: modules, data flow, hotspots, and dependency risks. Return Mermaid diagrams + a 90-day refactor plan. [link/tree]

Claude

ChatGPT

# NVIDIA NIM Prompt Catalog

Copy-paste prompts for build.nvidia.com "Try It!" boxes. Replace [brackets] with your values.

# LLMs (Chat Completions)

## 1) RAG summarizer — NIM: meta/llama-3.1-70b-instruct

```
System: Summarize with explicit citations like [S1], [S2]... User:
Question: [Q]. Sources: [S1 text] [S2 text] [S3 text]. Write 7 bullets;
each ends with citations. Finish with a one-line verdict.
```

## 2) OCR → tidy JSON tables — NIM: meta/llama-3.2-11b-vision-instruct

```
"Read every visible table and return only JSON: {tables:[{title,rows:
[{cells:[...]}]}]}. Preserve numbers/units; don't guess."
```

## 3) Chart interpreter — NIM: meta/llama-3.2-11b-vision-instruct

```
"Explain variables, trends, anomalies; give 3 causal hypotheses. End
with: Caution: [bias/limitation]."
```

## 4) Zero-shot class label — NIM: nvidia/nvclip

```
"a red vintage roadster in studio lighting; front-three-quarter view"
```

## 5) Fine-grained image query — NIM: nvidia/nv-dinov2

```
"close-up of [object] featuring [detail], background [style], captured
at [time of day]"
```

## 6) ASR Word boosting — NIM: nvidia/parakeet-ctc-1.1b-asr

```
"Boost list: AntiBERTa(40), ABlooper(40), your jargon, product.
Punctuation on; profanity filter on; numbers as digits."
```

## 7) ASR Fast English — NIM: nvidia/parakeet-tdt-0.6b-v2

```
"Transcribe with automatic punctuation + word timestamps. Prefer
standard casing."
```

## 8) Neutral explainer (SSML) — NIM: nvidia/magpie-tts-multilingual

```
<speaking> <prosody rate="92%" pitch="+1">Welcome to <sub alias="en vee
dee uh">NVIDIA</sub> NIM.</prosody> <break time="300ms"/> Today we'll
cover <emphasis level="moderate">three steps</emphasis>: setup, auth,
and integration. </speaking>
```

## 9) 5-second voice-match — NIM: nvidia/magpie-tts-flow

```
"Clone the attached 5-sec reference and read this: '[script]'. Target:
neutral corporate, 48kHz WAV, mono. Keep sibilance soft."
```

## 10) Lipsync from audio — NIM: nvidia/audio2face-3d

```
"Generate facial blendshapes from this narration. Output: ARKit-
compatible curves; neutralize head motion; smooth jaw at 6 Hz."
```

# 3D Asset Generation Pipeline

## A) Asset spec generator — NIM: nvidia/llama-3.1-nemotron-70b-instruct

```
System: You design production-ready asset specs. Output only the JSON
schema shown; no prose. User: "Create specs for [theme]. For each
asset, include: name, category (prop|environment|character),
description, dimensions_cm, poly_budget (low|mid|high), lods (e.g.,
LOD0, LOD1, LOD2), materials (PBR: albedo, normal, roughness,
metallic), file_format (gltf|usdz|usd), pivot (base|center),
scale_units (cm), licensing_note, ref_images (keywords). Return JSON: {
"assets": [ {...} ], "theme": "[theme]" }."
```

## B) 3D generation — NIM: microsoft/trellis

```
"Generate a physically coherent [category] named [name] matching this
description: [description]. Constraints: dimensions ≈ [dimensions_cm]
cm, [poly_budget] polycount, provide [lods], PBR textures
(albedo/normal/roughness/metallic), export as [file_format] with real-
world cm scale, pivot at [pivot]. Ensure watertight meshes; no non-
```

```
manifold edges; correct face normals; clean UVs; texture resolution ≥ 2K."
```

### C) Audio generation — NIM: `nvidia/magpie-tts-zeroshot`

```
"Clone the attached 5-sec sample and generate [N] lines for the asset pack '[theme]'. Tone [tone], language [lang], sample rate 48 kHz, mono WAV, loudness -16 LUFS. Lines: [bulleted lines]. Return a file list with durations."
```

### D) Validation — NIM: `nvidia/usdvalidate`

```
"Validate these assets for OpenUSD. Enforce: cm units; pivot at [pivot]; outward normals; no non-manifold geometry; no inverted UVs; texture existence and ≥ 2048px; LOD naming LOD0/LOD1/LOD2; axis-aligned bounds; material-slot count matches spec. Return JSON report: {file, errors[], warnings[], metrics:{tris, materials, texResMin}}."
```

### E) If-fail auto-regen — NIM: `nvidia/llama-3.1-nemotron-70b-instruct`

```
"Given this validator report [paste JSON], propose minimal constraint tweaks per asset (e.g., reduce poly, fix UV seams, raise texture res). Output JSON patch: {name, changes[], rationale}. Keep changes few and surgical."
```

### F) Manifest + QA roll-up — NIM: `nvidia/llama-3.1-nemotron-70b-instruct`

```
"Combine the final asset specs and validator reports into a single manifest: {assets:[{name, file, dims_cm, tris, lods, materials, texResMin, issues_fixed[]}], theme, created_at}. Produce a 5-bullet QA summary."
```



## Key Takeaways & Best Practices



## Iteration is Key

The first prompt is

